

Prague AWS User Group

Running Docker Containers on AWS (with ECS)

Vladimir Simek, Solutions Architect @ AWS
20 September 2017

Agenda

Why Containers?

Cluster Management

Running Services

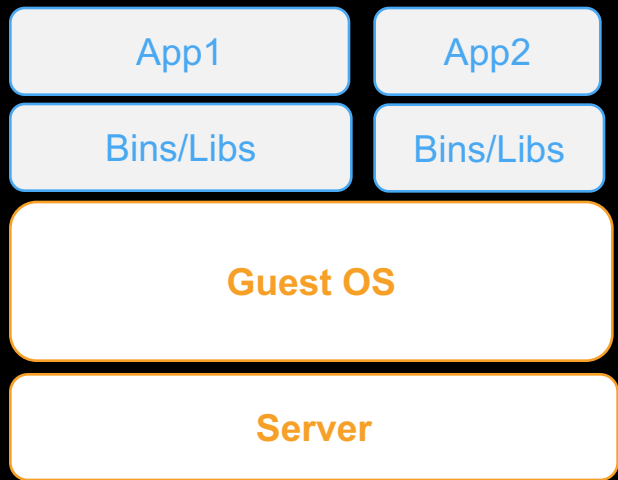
Placement Options

Security and Monitoring

Demo

Why Containers?

What are Containers?



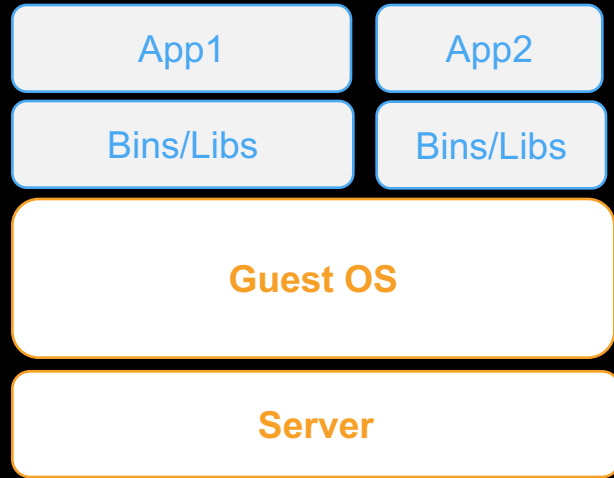
OS virtualization

Process isolation

Images

Automation

Container advantages



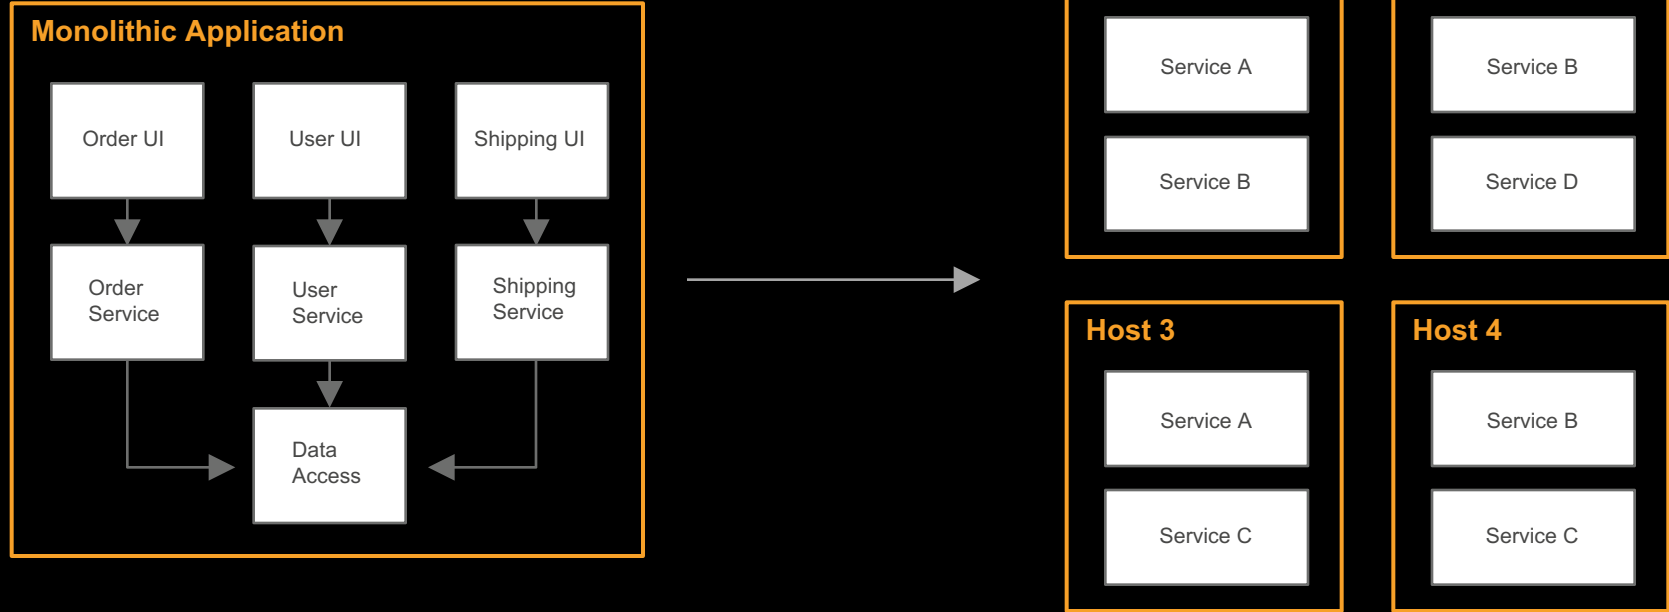
Portable

Flexible

Fast

Efficient

Services evolve to microservices



Containers are natural for microservices

Simple to model

Any app, any language

Image is the version

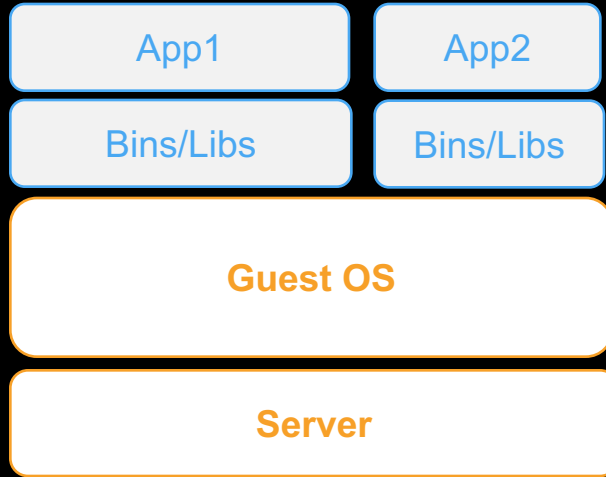
Test & deploy same artifact

Stateless servers decrease change risk


Scheduling



Scheduling one resource is straightforward



Scheduling a cluster is hard

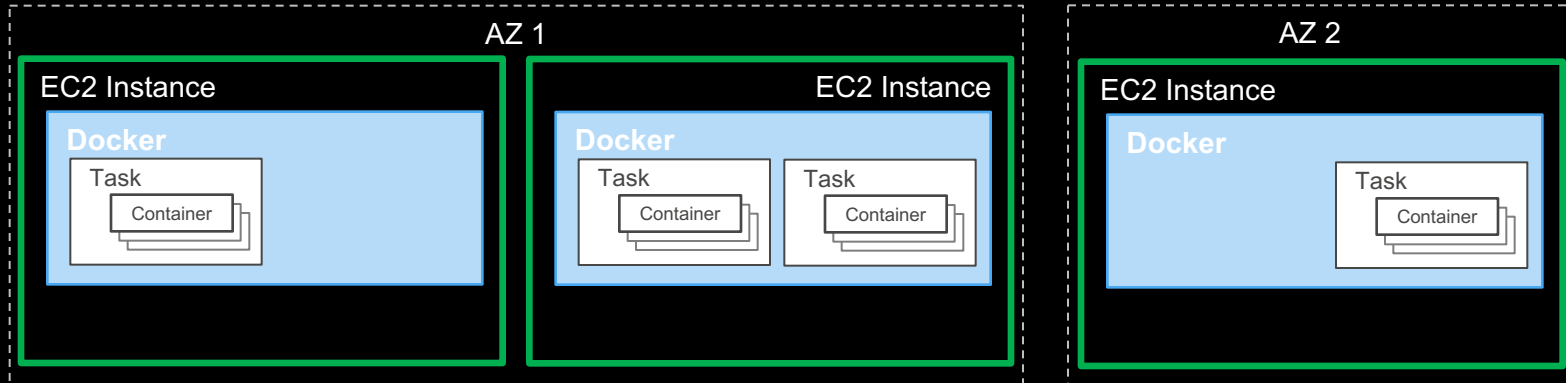


What is Amazon ECS?

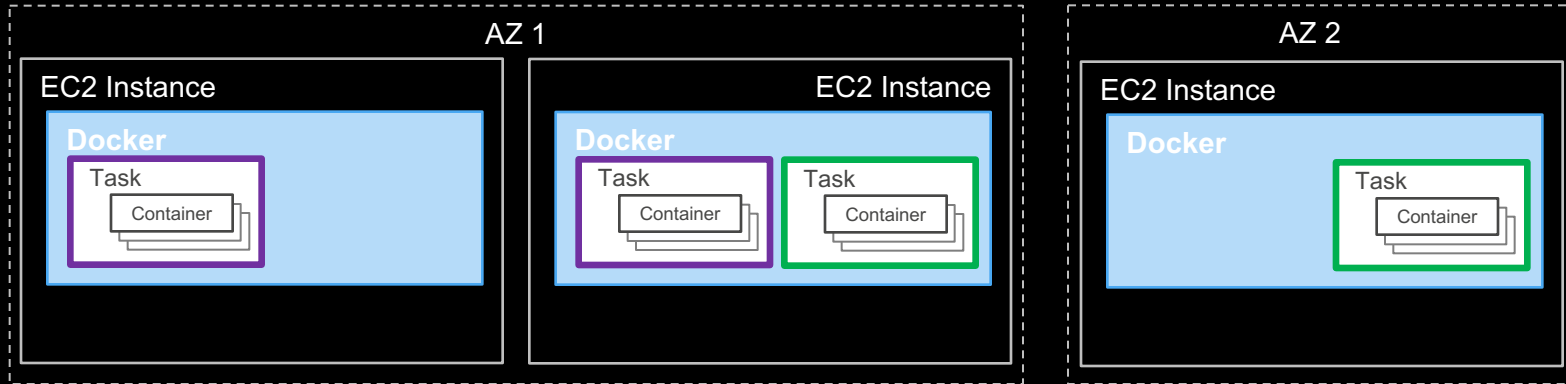
Amazon EC2 Container Service (ECS) is a highly scalable, high performance **container management service**. You can use Amazon ECS to **schedule** the placement of containers across your cluster. You can also integrate your own **scheduler** or **third-party scheduler** to meet business or application specific requirements.

Cluster Management

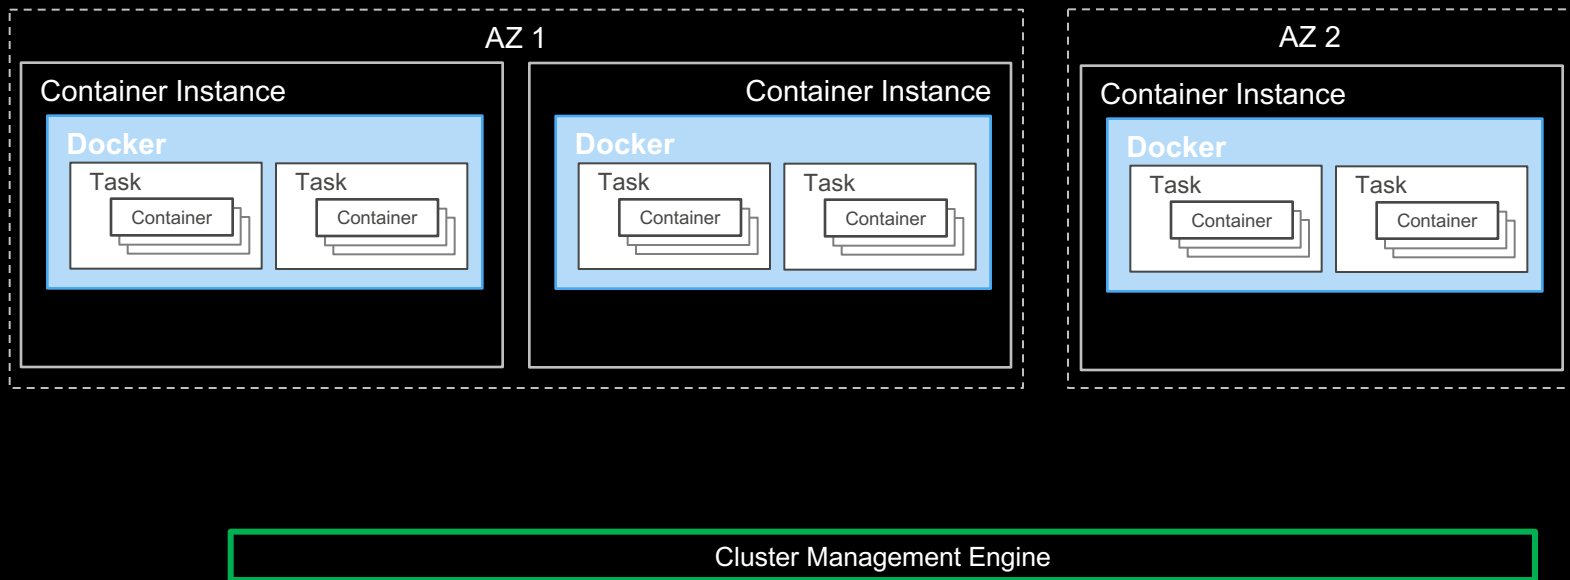
Cluster Management: Resource Management



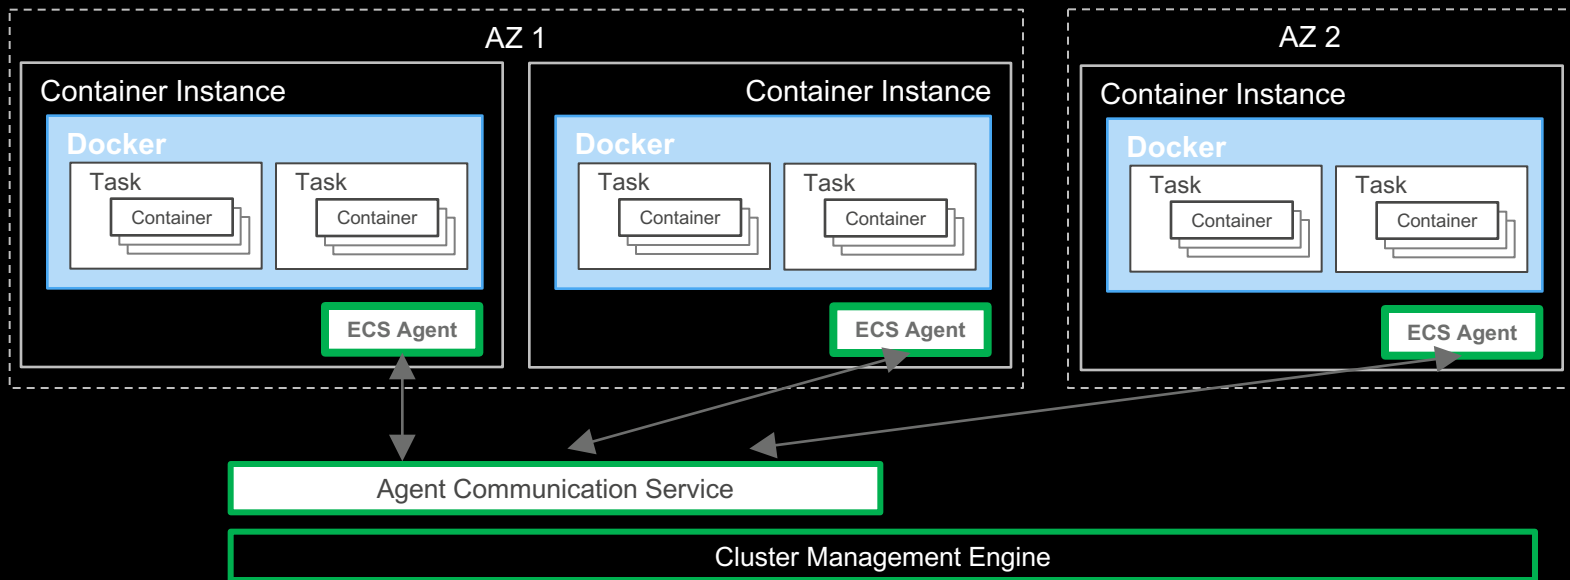
Cluster Management: Scheduling



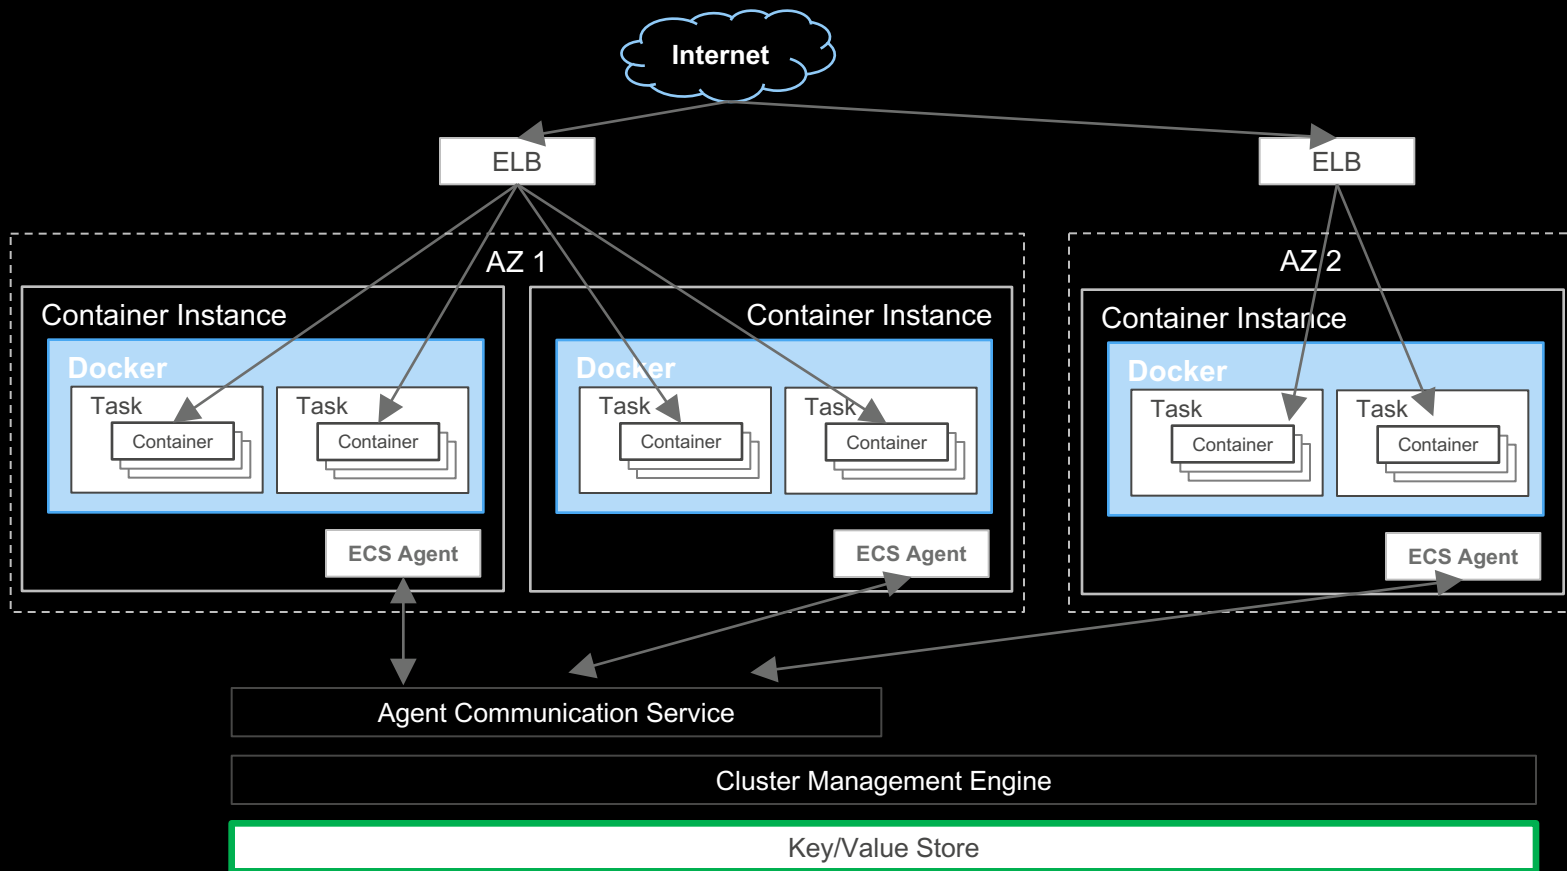
Amazon ECS: Resource Management



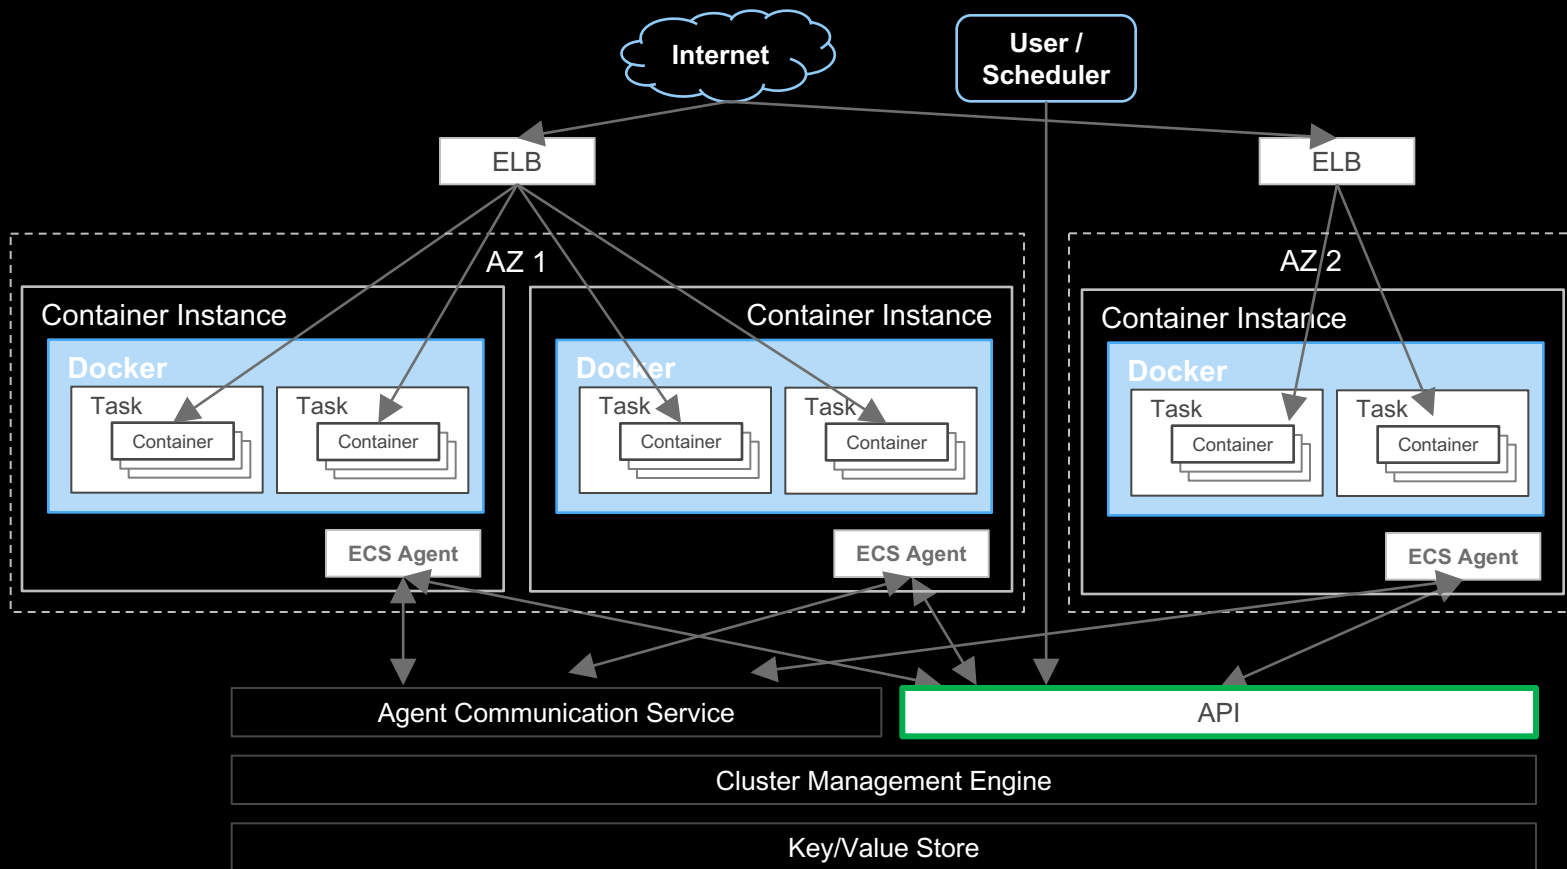
Amazon ECS: Agent Communication



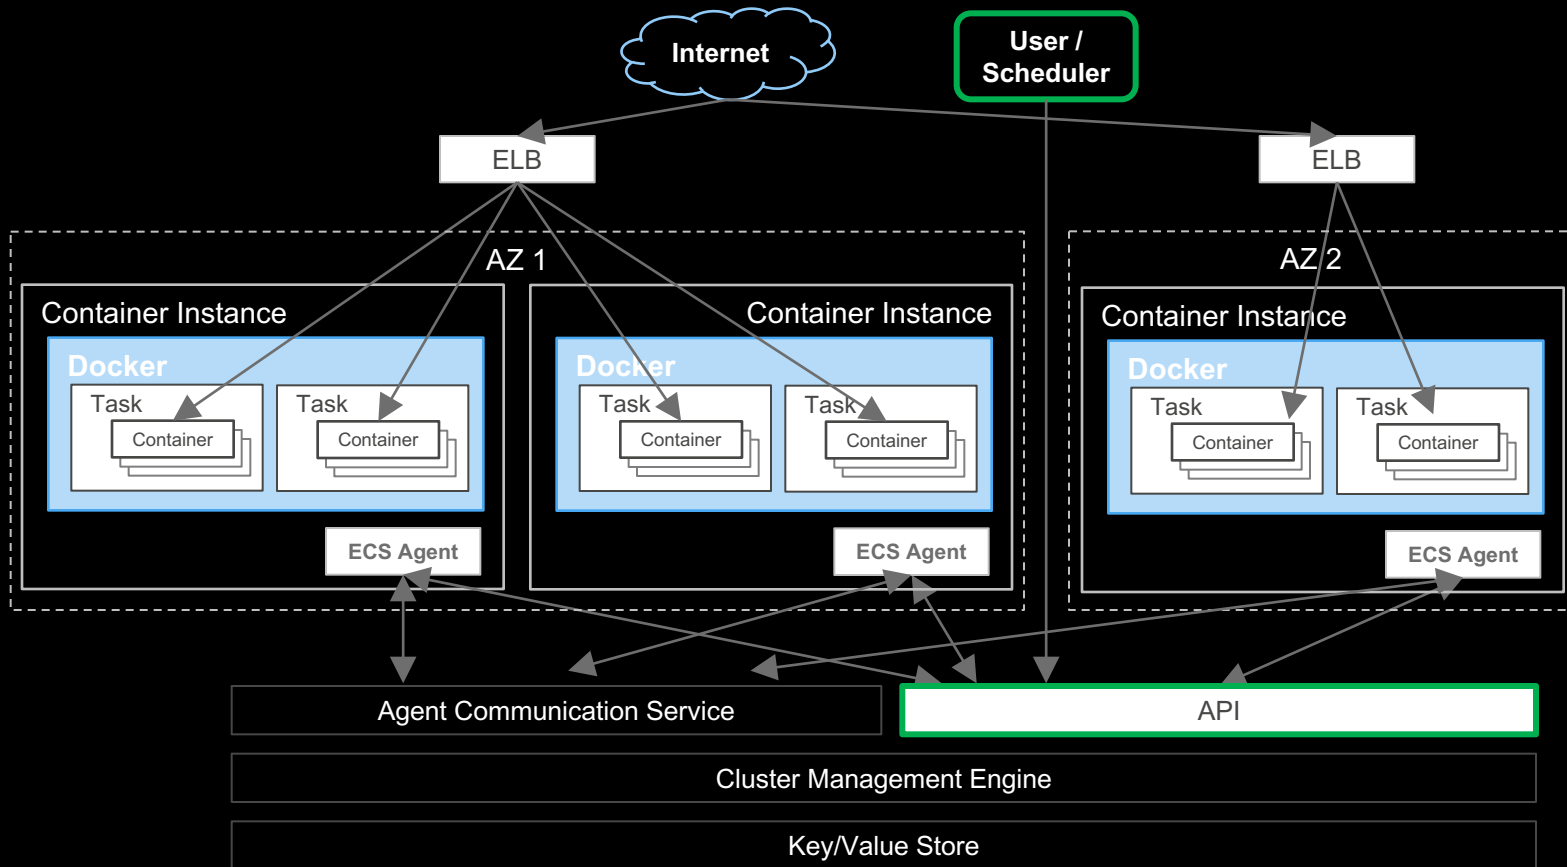
Amazon ECS: Key/Value Store



Amazon ECS: APIs




Amazon ECS: Scheduling



Benefits

Easily Manage Clusters for Any Scale




Nothing to run

Complete state


Control and monitoring

Scale

Scalable



Flexible Container Placement




Applications

Batch jobs

Multiple schedulers

Designed for use with other AWS services



Elastic Load Balancing

Amazon Elastic Block Store

Amazon Virtual Private Cloud

Amazon CloudWatch

AWS Identity and Access Management

AWS CloudTrail

Extensible



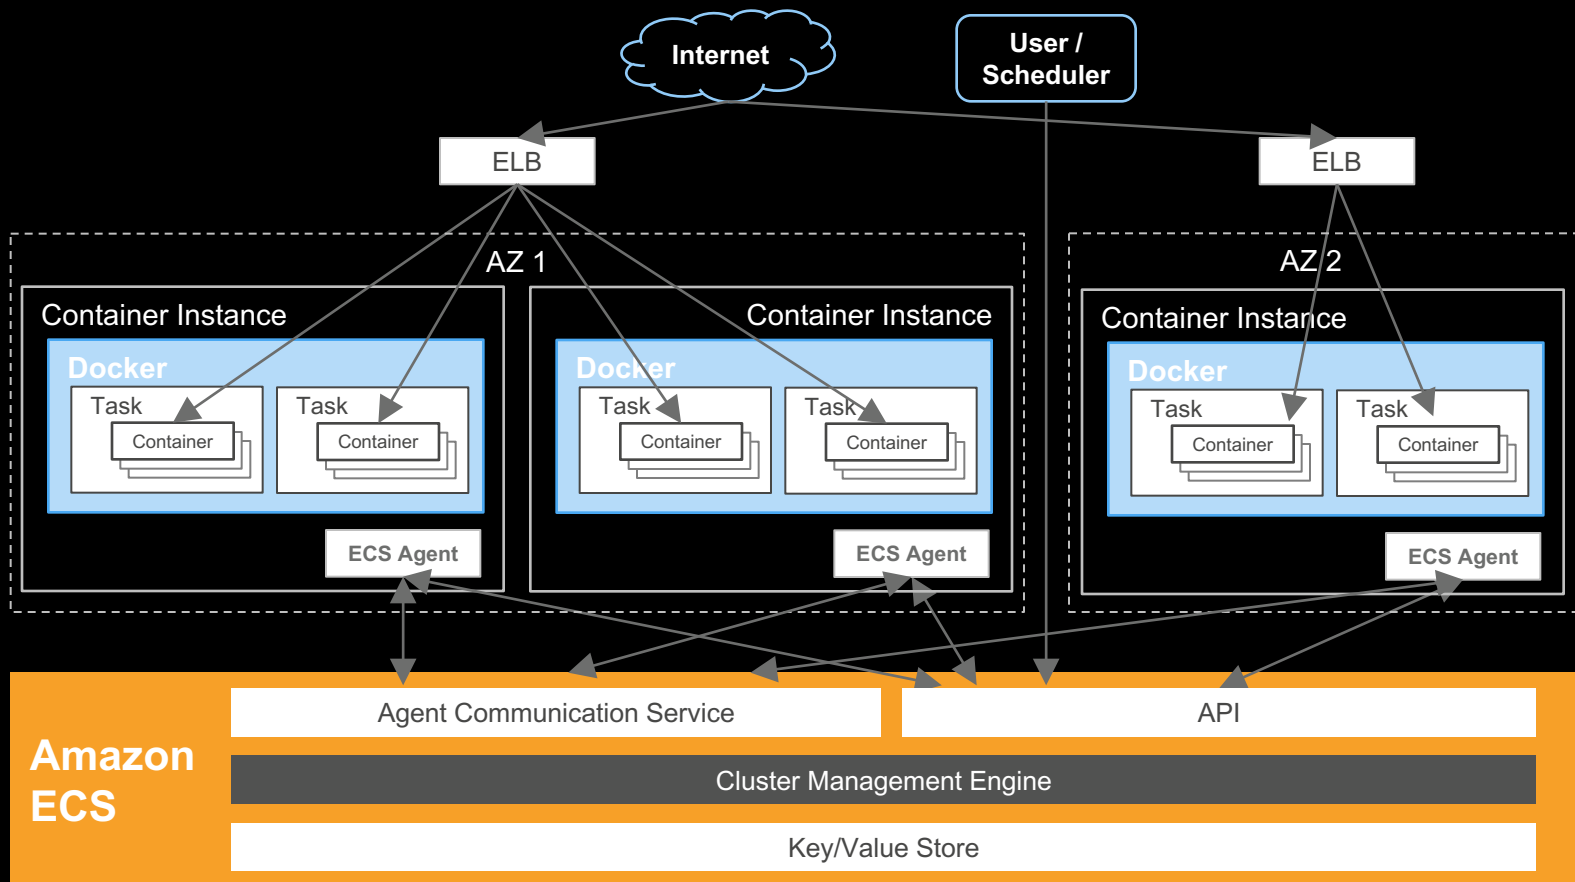
Comprehensive APIs

Custom schedulers

Open source agent and CLI

Blox

Amazon ECS



ECS Workflow

Typical User Workflow



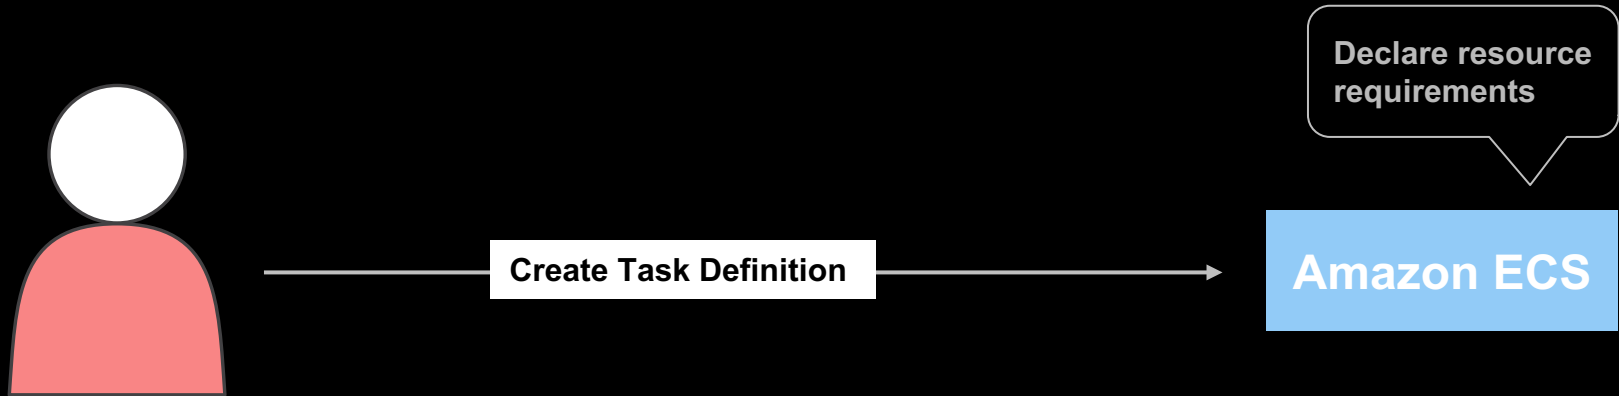
Typical User Workflow



Push Image to ECR

```
$ docker push my-account.dkr.ecr.us-east-1.amazonaws.com/myImage
```

Typical User Workflow



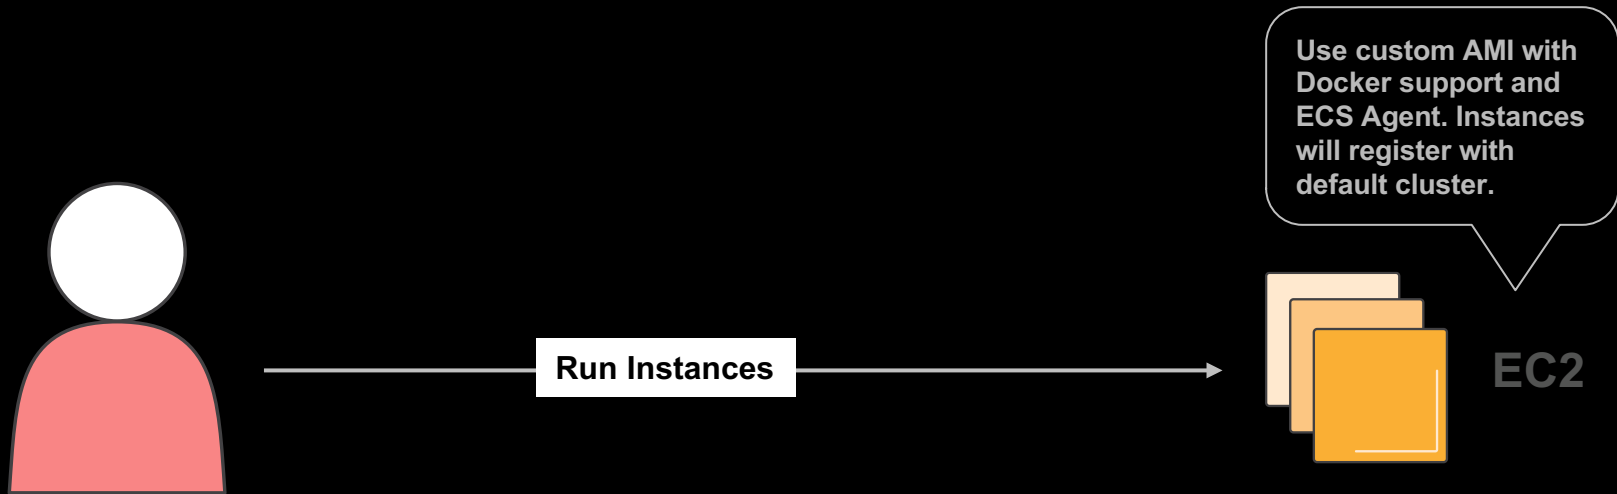
Register task definition

```
$ aws ecs register-task-definition --cli-  
input-json file://task-defintion.json
```


Task definition

```
{
  "taskDefinition": {
    "status": "INACTIVE",
    "family": "curler",
    "volumes": [],
    "taskDefinitionArn": "arn:aws:ecs:us-west-2:123456778901:task-definition/curler:1",
    "containerDefinitions": [
      {
        "environment": [],
        "name": "curler",
        "mountPoints": [],
        "image": "curl:latest",
        "cpu": 100,
        "portMappings": [],
        "entryPoint": [],
        "memory": 256,
        "command": [
          "curl -v http://example.com/"
        ],
        "essential": true,
        "volumesFrom": []
      }
    ],
    "revision": 1
  }
}
```

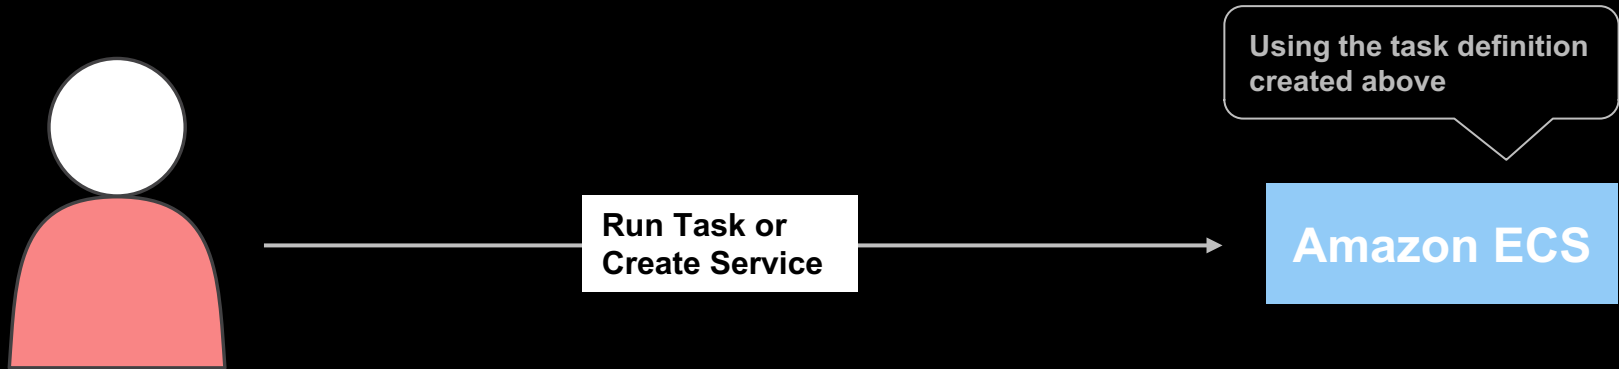
Typical User Workflow



Create cluster

```
$ aws ecs create-cluster --cluster-name  
"my_cluster"
```

Typical User Workflow



Create service

```
$ aws ecs create-service --service-name ecs-simple-service  
--task-definition ecs-demo --desired-count 10
```


Create Service

Load Balance traffic across containers

Automatically recover unhealthy containers

Discover services

Elastic Load Balancing







Scale Service

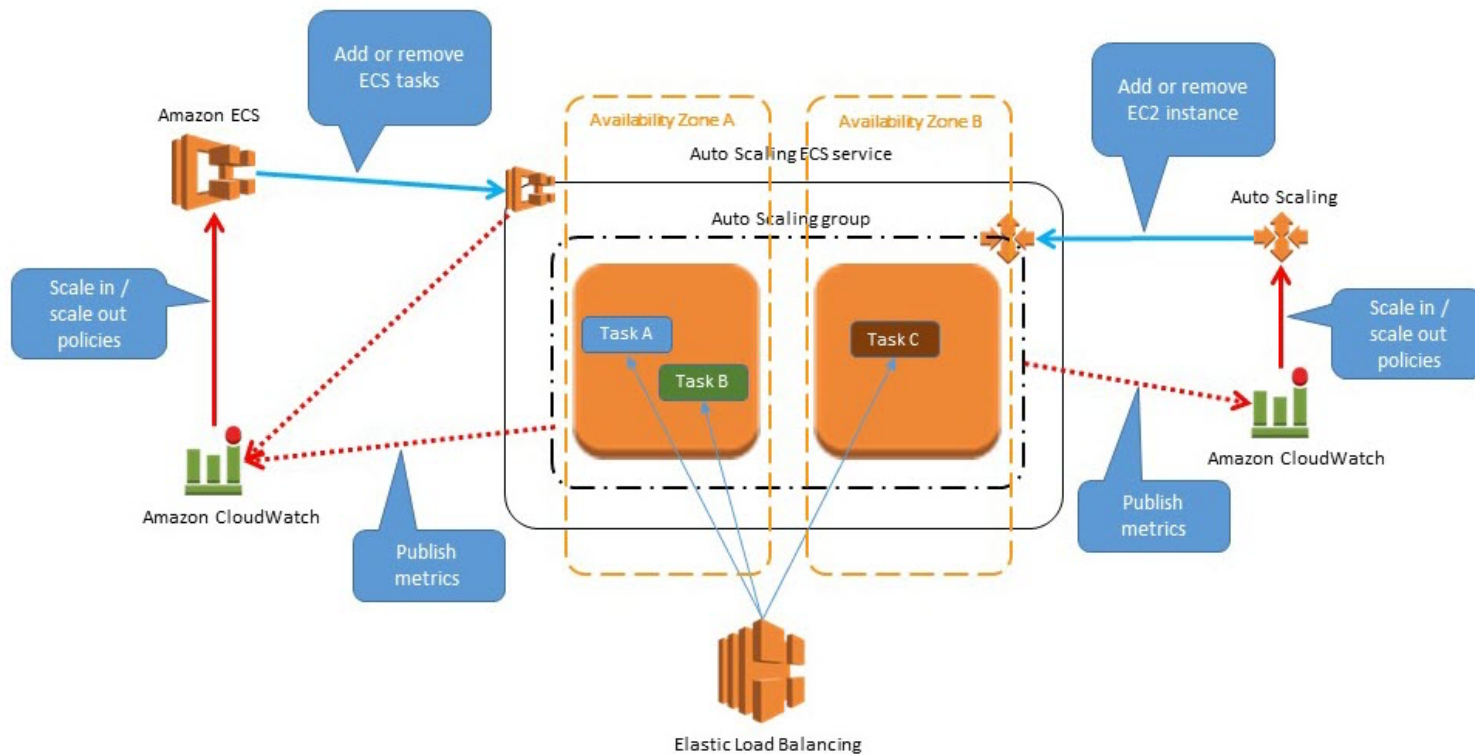
Scale out

Scale in

Elastic Load Balancing



Scale Service

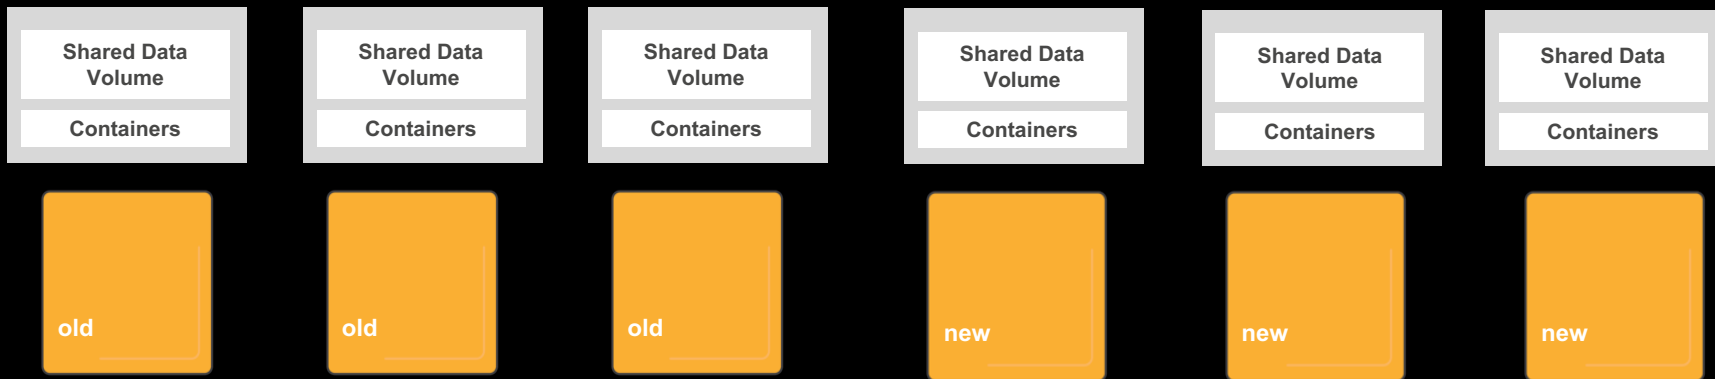


Update Service

Deploy new version

Drain connections

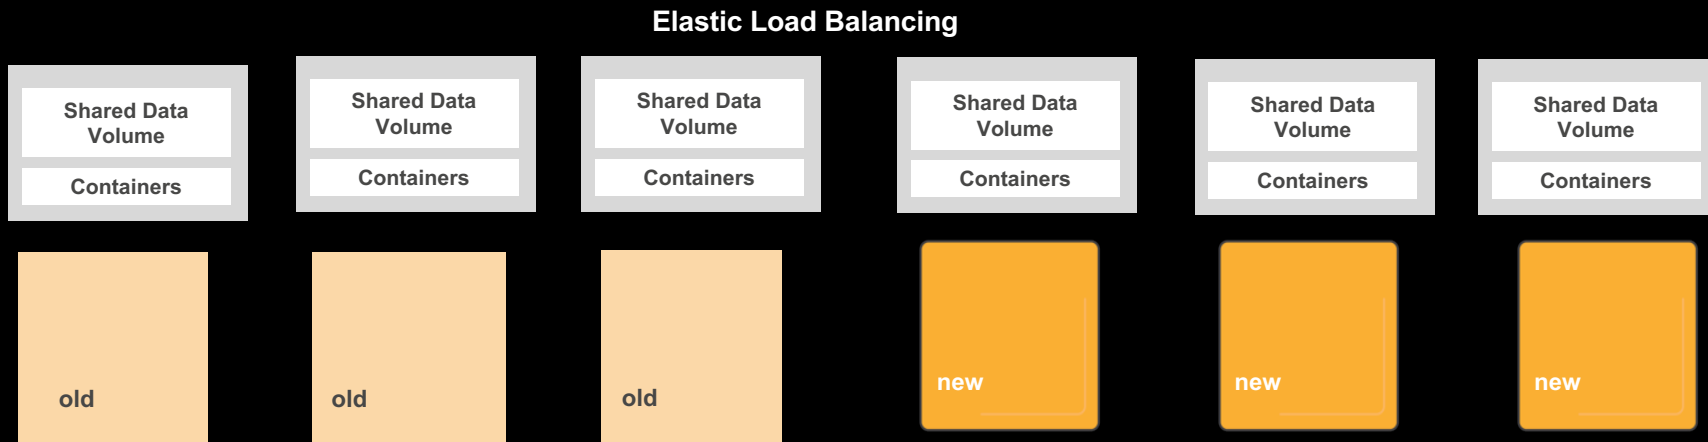
Elastic Load Balancing



Update Service (cont.)

Deploy new version

Drain connections




Update Service (cont.)

Deploy new version

Drain connections

Elastic Load Balancing



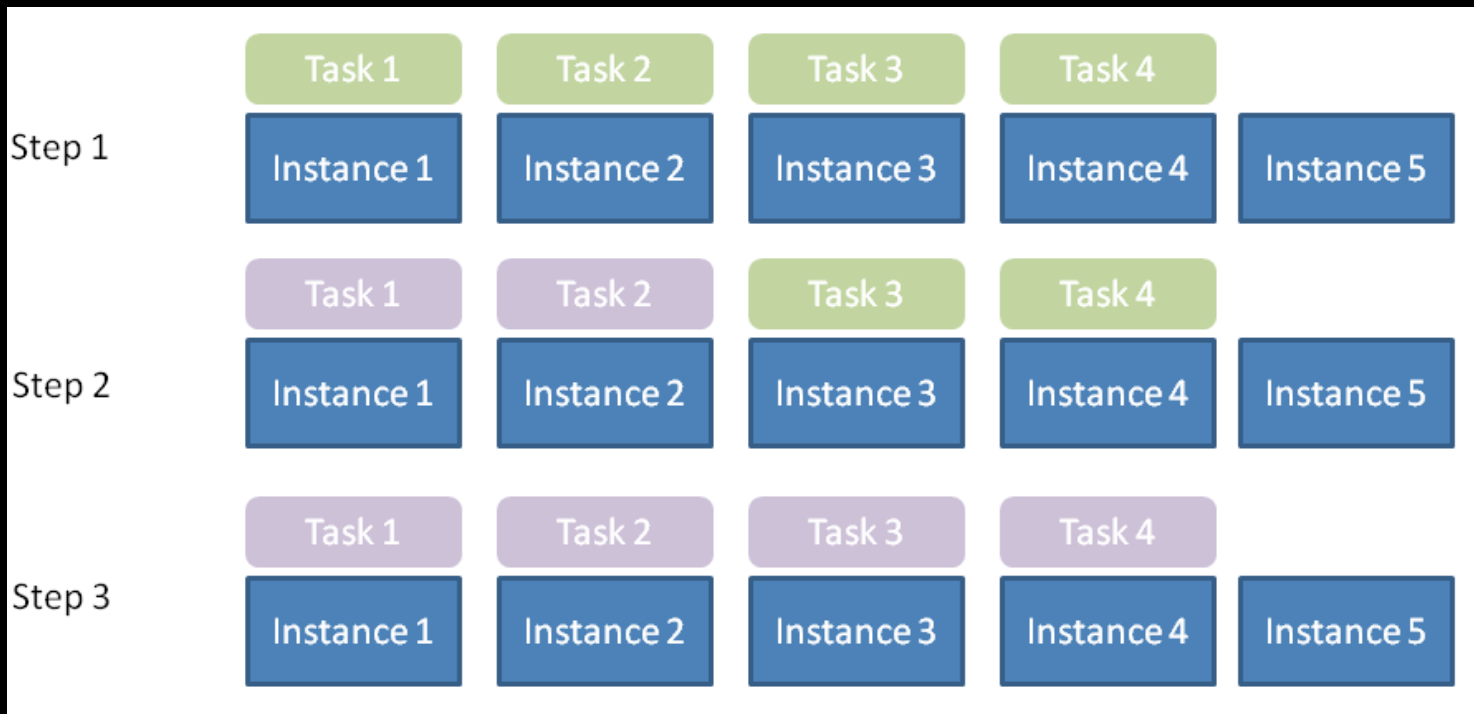
Update Service (cont.)

Specify a deployment configuration for your service:

- *minimumHealthyPercent*: lower limit (as a percentage of the service's desiredCount) of the number of running tasks that must remain running in a service during a deployment.
- *maximumPercent*: upper limit (as a percentage of the service's desiredCount) of the number of running tasks that can be running in a service during a deployment.

Update Service (cont.)

Deploy using the least space: *minimumHealthyPercent* = 50%, *maximumPercent* = 100%




Update Service (cont.)

Deploy quickly without reducing service capacity:
minimumHealthyPercent = 100%, *maximumPercent* = 200%




Task Placement Examples


Supported Placement Strategies




Binpacking



Spread




Affinity



Distinct Instance

Placement: Targeting Instance Type

```
aws ecs run-task --cluster ecs-demo --task-definition myapp --count 5 --placement-constraints  
type="memberOf",expression="attribute:ecs.instance-type == g2.2xlarge"
```



g2.2xlarge



t2.small




g2.2xlarge



g2.2xlarge



Placement: Targeting Instance Type & Zone

```
aws ecs run-task --cluster ecs-demo --task-definition myapp --count 5 --placement-constraints  
type="memberOf",expression="(attribute:ecs.instance-type == t2.small or  
attribute:ecs.instance-type == t2.medium) and attribute:ecs.availability-zone != us-east-1d"
```



Placement: Availability Zone Spread

```
aws ecs run-task --cluster ecs-demo --task-definition myapp --count 9 --placement-strategy  
type="spread",field="attribute:ecs.availability-zone"
```



Security

IAM Roles

IAM roles for **container instances**:

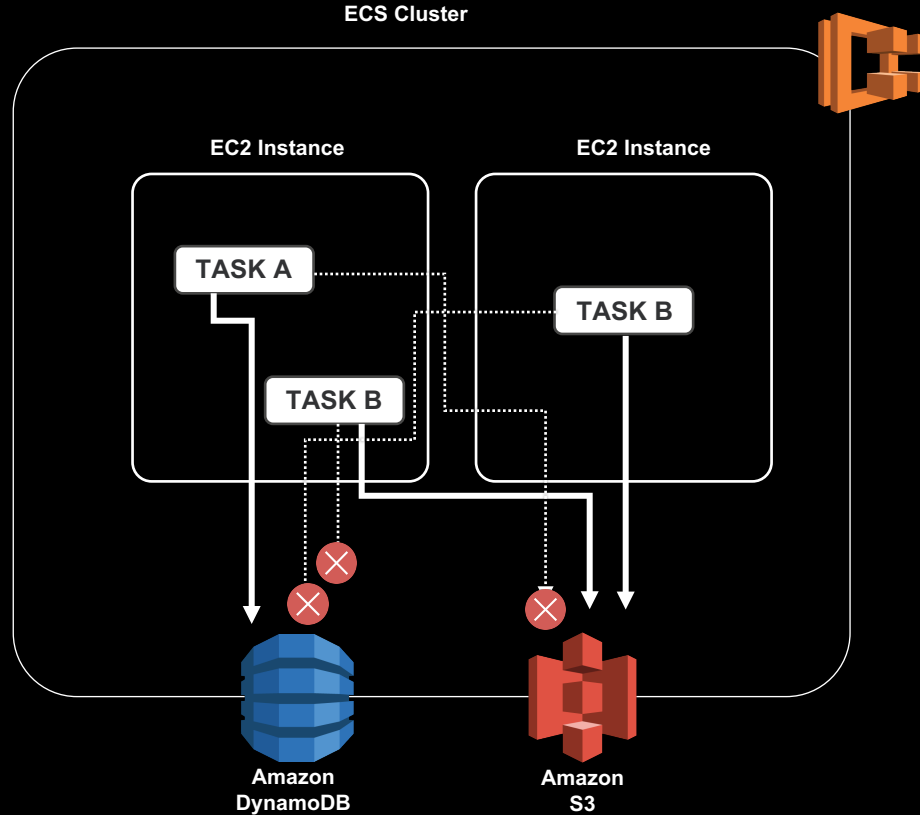
- Bound to the ECS container instance
- Applies to all containers running on the host
- Pulling images from ECR
- CloudWatch Logs

IAM roles for **tasks**:

- Bound to specific ECS tasks
- Task-specific access to AWS services

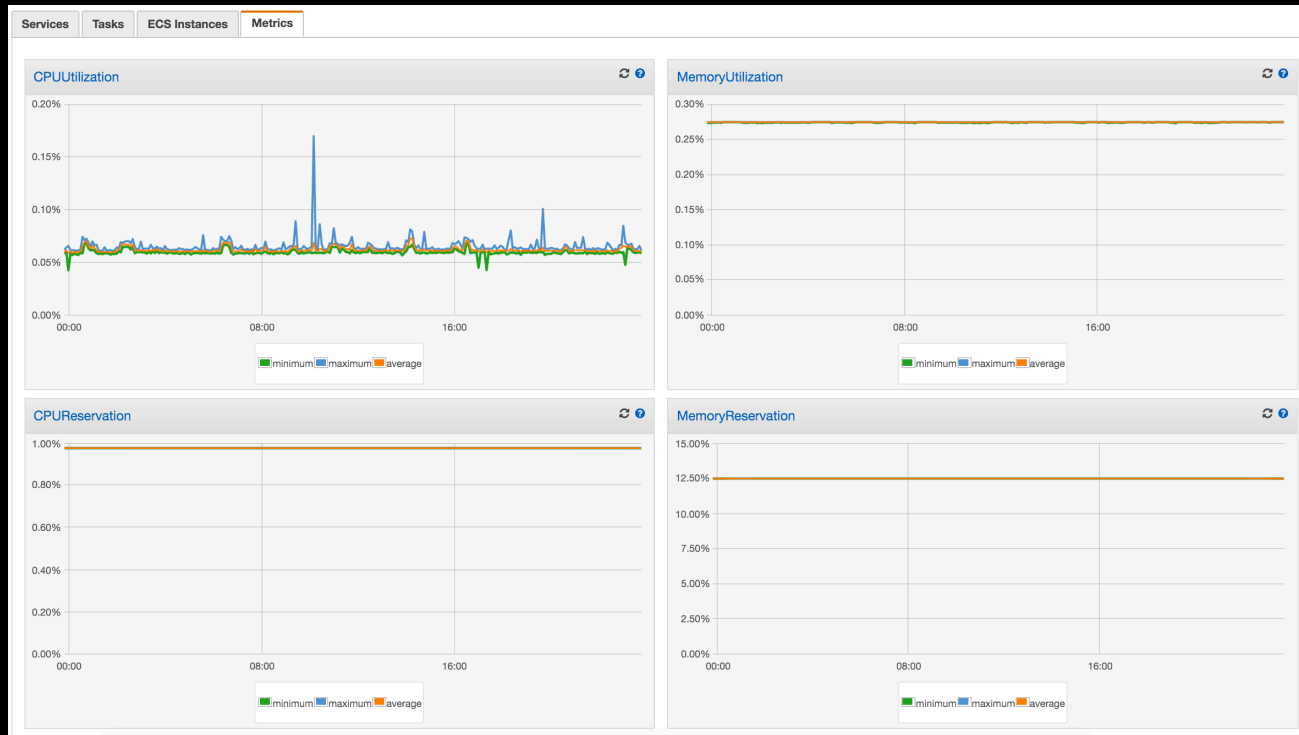
Tip Use principle of least privilege – prefer **IAM roles for tasks** where applicable

IAM Roles For Tasks



Monitoring & Logging

Monitoring with CloudWatch



Centralized Logging with CloudWatch Logs

```
{
  "image": "nginx:latest",
  ...
  "logConfiguration": {
    "logDriver": "awslogs",
    "options": {
      "awslogs-group": "nginx",
      "awslogs-region": "us-east-1"
    }
  }
}
```

- Defined within the task definition
- Available log drivers
 - awslogs
 - fluentd
 - gelf
 - journald
 - json-file
 - splunk
 - Syslog
- Submit a pull request on ECS agent GitHub repo if you would like others

Centralized Logging with CloudWatch Logs

CloudWatch > Log Groups > Production-WebsiteService > 8a2d684337b2cf37c324f221c697ea0c28d8bd841a40a709a540b8fc7957a360

Expand all ☒ Row ☐ Text



200



all 30s 5m 1h 6h 1d 1w custom ▾

	Time (UTC +00:00)	Message
2016-11-03		
▶	21:48:59	[GIN] 2016/11/03 - 21:48:59 [[97;42m 200 [0m] 20.560916ms 10.180.23.86:1605 [[97;44m [0m GET /
▶	21:48:59	[GIN] 2016/11/03 - 21:48:59 [[97;42m 200 [0m] 20.170986ms 10.180.14.27:63074 [[97;44m [0m GET /
▶	21:49:09	[GIN] 2016/11/03 - 21:49:09 [[97;42m 200 [0m] 3.23317ms 10.180.23.86:1625 [[97;44m [0m GET /
▶	21:49:09	[GIN] 2016/11/03 - 21:49:09 [[97;42m 200 [0m] 853.79µs 10.180.14.27:63096 [[97;44m [0m GET /
▶	21:49:19	[GIN] 2016/11/03 - 21:49:19 [[97;42m 200 [0m] 918.472µs 10.180.23.86:1633 [[97;44m [0m GET /
▶	21:49:19	[GIN] 2016/11/03 - 21:49:19 [[97;42m 200 [0m] 879.181µs 10.180.14.27:63102 [[97;44m [0m GET /
▶	21:49:29	[GIN] 2016/11/03 - 21:49:29 [[97;42m 200 [0m] 1.113159ms 10.180.23.86:1643 [[97;44m [0m GET /
▶	21:49:29	[GIN] 2016/11/03 - 21:49:29 [[97;42m 200 [0m] 921.021µs 10.180.14.27:63110 [[97;44m [0m GET /
▶	21:49:39	[GIN] 2016/11/03 - 21:49:39 [[97;42m 200 [0m] 2.395537ms 10.180.23.86:1653 [[97;44m [0m GET /
▶	21:49:39	[GIN] 2016/11/03 - 21:49:39 [[97;42m 200 [0m] 971.799µs 10.180.14.27:63118 [[97;44m [0m GET /
▶	21:49:49	[GIN] 2016/11/03 - 21:49:49 [[97;42m 200 [0m] 1.163781ms 10.180.23.86:1661 [[97;44m [0m GET /
▶	21:49:49	[GIN] 2016/11/03 - 21:49:49 [[97;42m 200 [0m] 970.934µs 10.180.14.27:63128 [[97;44m [0m GET /
▶	21:49:59	[GIN] 2016/11/03 - 21:49:59 [[97;42m 200 [0m] 1.10307ms 10.180.23.86:1667 [[97;44m [0m GET /
▶	21:49:59	[GIN] 2016/11/03 - 21:49:59 [[97;42m 200 [0m] 870.338µs 10.180.14.27:63136 [[97;44m [0m GET /
▶	21:50:09	[GIN] 2016/11/03 - 21:50:09 [[97;42m 200 [0m] 1.058436ms 10.180.23.86:1681 [[97;44m [0m GET /
▶	21:50:09	[GIN] 2016/11/03 - 21:50:09 [[97;42m 200 [0m] 898.432µs 10.180.14.27:63148 [[97;44m [0m GET /

Demo

AWS Customers using ECS in production

airtime

Atlassian

Aol.

Bb
Blackboard

edmunds.com

Expedia

F
FOURSQUARE

GILT

here

instacart

lyft

meetup

X mytaxi

okta

RAC
Retail & Commerce
FURNITURE • APPLIANCES • ELECTRONICS
COMPUTERS

Riot
GAMES

slack

SUP
ERC
ELL

Trimble

Upserve



14

Data Processing
Services

3500

Peak Container
Instances

500 million

Compute Hours Used
This Year



- 50 Microservices - 10 containers/service
- Costs down 40% by using Spot Instances
- Faster Deployment Times

“In November 2015 we moved our Docker container architecture to [Amazon ECS](#), and for the first time ever in December we were able to celebrate a new year in which our system could handle the huge number of requests without any crashes or interruptions—an accomplishment that we were extremely proud of. We had faced the biggest night on the calendar without any downtime.”


Docker on AWS – just ECS?



kubernetes



MESOS



AWS & CNCF



Questions?

Thank you

vladsim@amazon.com